


Dr. Alex Johnson

(+1) 555-123-4567 | 456 Innovation Drive, San Francisco, CA 94105 | alex.johnson@techcorp.com | alexjohnson-ml | alexjohnson.tech | <https://alexjohnson.tech> |  orcid.org/0000-0002-1234-5678

Education

Stanford University

M.S. in Computer Science (Machine Learning)

GPA: 3.9/4.0 • Honors: Graduated Summa Cum Laude

Relevant Coursework: Deep Learning, Reinforcement Learning, Computer Vision, Natural Language Processing

- Thesis: "Scalable Federated Learning for Edge Computing" (Advisor: Prof. Andrew Ng)
- Teaching Assistant for CS229 (Machine Learning) and CS231n (Convolutional Neural Networks)

August 2015 - May 2017

Stanford, CA

Massachusetts Institute of Technology

B.S. in Electrical Engineering and Computer Science

GPA: 3.8/4.0 • Honors: Dean's List (2012-2015), Phi Beta Kappa Honor Society

Relevant Coursework: Algorithms, Artificial Intelligence, Signal Processing, Linear Algebra

- Senior Thesis: "Real-time Object Detection for Autonomous Vehicles"
- President of MIT AI/ML Club • Organized annual hackathon with 500+ participants

August 2011 - May 2015

Cambridge, MA

Experience

TechCorp Industries

Senior Machine Learning Engineer

- Architected and deployed scalable ML pipelines processing 10M+ daily transactions with 99.9% uptime
- Led cross-functional team of 8 engineers to develop real-time recommendation system, increasing user engagement by 35%
- Implemented MLOps practices reducing model deployment time from weeks to hours using Kubernetes and Docker
- Mentored 12 junior engineers and established ML engineering best practices across the organization

2021 - Present

San Francisco, CA

AI Innovations Lab

Research Scientist

- Pioneered novel deep learning architectures for computer vision achieving state-of-the-art results on ImageNet
- Published 8 peer-reviewed papers in top-tier conferences (NeurIPS, ICML, ICLR) with 200+ total citations
- Secured \$2.5M in research funding through successful grant proposals to NSF and industry partnerships
- Collaborated with Stanford and MIT researchers on breakthrough federated learning algorithms

2019 - 2021

Mountain View, CA

StartupTech, Inc.

Machine Learning Engineer

- Developed end-to-end ML solutions for fraud detection reducing false positives by 60% and saving \$5M annually
- Built distributed training infrastructure supporting models with 100B+ parameters using PyTorch and Ray
- Optimized inference pipeline achieving 10x latency improvement through model quantization and hardware acceleration

2017 - 2019

Palo Alto, CA

Projects

AutoML Platform

[alexjohnson-ml/automl-platform](#)

- Designed comprehensive AutoML platform enabling non-technical users to build ML models with 90% accuracy
- Implemented distributed hyperparameter optimization reducing training time by 75% using Optuna and Ray
- Platform adopted by 500+ internal users across 15 business units, accelerating time-to-insight by 80%
- Built web interface using React and FastAPI with real-time model monitoring and A/B testing capabilities

Jan 2023 - Present

Federated Learning Framework

[alexjohnson-ml/federated-ml](#)

- Developed privacy-preserving federated learning framework supporting 1000+ edge devices
- Implemented differential privacy mechanisms ensuring ϵ -differential privacy with minimal accuracy loss
- Framework enables collaborative ML training across organizations without data sharing
- Open-sourced project with 2.5K GitHub stars and active contributor community

Mar 2022 - Dec 2022

Real-time Anomaly Detection System

[alexjohnson-ml/anomaly-detection](#)

- Built production-grade anomaly detection system processing 50M events/day with <100ms latency
- Utilized ensemble of LSTM and Transformer models achieving 95% precision and 92% recall
- Integrated with Kafka, ClickHouse, and Grafana for real-time monitoring and alerting
- System detected critical infrastructure issues 30 minutes faster than previous solutions

Aug 2021 - Feb 2022

Skills

Programming Languages

Python, C++, JavaScript, Go, Rust, Java, R, CUDA

Machine Learning

PyTorch, TensorFlow, Scikit-learn, XGBoost, Hugging Face, MLflow, Ray, Optuna

Cloud & Infrastructure

AWS, Kubernetes, Docker, GCP, Apache Kafka, Redis, PostgreSQL, ClickHouse

Specialized Skills

Deep Learning, Computer Vision, NLP, Federated Learning, MLOps, A/B Testing

Publications & Research

Federated Learning with Differential Privacy: A Comprehensive Survey

2024

Nature Machine Intelligence

First Author • Impact Factor: 25.9 • Cited by 156

Efficient AutoML for Large-Scale Industrial Applications

2023

International Conference on Machine Learning (ICML)

Lead Author • Acceptance Rate: 27.9 • Cited by 89

Privacy-Preserving Deep Learning at Scale

2022

Neural Information Processing Systems (NeurIPS)

Co-Author • Top 1% of submissions • Cited by 234

Certifications & Awards

AWS Certified Machine Learning - Specialty

2023

Google Cloud Professional ML Engineer

2022

Certified Kubernetes Administrator (CKA)

2023

- **Best Paper Award** - International Conference on Machine Learning Applications (2023)
- **Rising Star in AI Award** - AI Innovation Summit (2022)
- **Outstanding Graduate Award** - Stanford University (2017)
- **NSF Graduate Research Fellowship** (2015-2017)